



developerWorks   Technical topics   Rational   Technical library

# UML basics: The class diagram

## An introduction to structure diagrams in UML 2

from The Rational Edge: As the most important example of the new structure diagram, this diagram can be used by analysts, business modelers, developers, and testers through lifecycle. This article offers a comprehensive introduction.

Donald Bell is an IT architect for IBM Rational software, where he works with clients to define and adopt and tools.



present in a series of articles

diagrams used within the

page, or UML. In my

previous article on [sequence diagrams](#), I shifted

focus away from the UML 1.4 spec to OMG's

ion of UML (a.k.a.

discuss Structure Diagrams, which is a new

been introduced in UML 2. Because the purpose

ements and their meanings, this article focuses

become clear. Subsequent articles will cover

this is the next

structure category.

installment in a series of articles about the

essential to reading and understanding this series is that UML notation elements

Modeling to provide guidance on the best approach for modeling, or how

should be modeled in the first place. Instead, the purpose of this article ar

on "sequence diagrams", I shifted focus away



*help with a basic understanding of notation elements — their syntax and the knowledge you should be able to read diagrams and create your own diagrams.*

*This article assumes you have a rudimentary understanding of object-oriented programming. If you need a little assistance with OO concepts, you might try the Sun brief [Programming Concepts](#). Reading the sections "What Is a Class?" and "What You Need to Know to Make This Article Useful." In addition, David H.arel Technologies: A Manager's Guide, offers an excellent, high-level explanation without requiring an in-depth understanding of computer programming.*

## The yin and yang of UML 2

In UML 2 there are two basic categories of diagrams: structure diagrams and behavior diagrams. Every UML diagram belongs to one of these two diagram categories. The purpose of structure diagrams is to show the static structure of the system being modeled. They include the class, component, and object diagrams. Behavioral diagrams, on the other hand, show the dynamic behavior between the objects in the system, including things like their methods, collaborations, and activities. Example behavior diagrams are activity, use case, and sequence diagrams.

Re  
E  
U  
U  
n  
V  
IE  
C  
th  
V  
IE  
C  
B

## Structure diagrams in general

As I have said, structure diagrams show the static structure of the system elements of a system, irrespective of time. Static structure is conveyed by instances in the system. Besides showing system types and their instances, they show at least some of the relationships among and between these elements and their internal structure.

Structure diagrams are useful throughout the software lifecycle for a variety of purposes. These diagrams allow for design validation and design communication between stakeholders. For example, business analysts can use class or object diagrams to model business entities and resources, such as account ledgers, products, or geographic hierarchies. Developers can use component and deployment diagrams to test/verify that their design is sound. Architects can use diagrams to design and document the system's coded (or soon-to-be-coded) components.

## The class diagram in particular

UML 2 considers structure diagrams as a classification; there is no diagram type called "Class Diagram." However, the class diagram offers a prime example of the structure diagram. It provides us with an initial set of notation elements that all other structure diagrams build upon. Since the class diagram is so foundational, the remainder of this article will focus on it. By the end of this article you should have an understanding of how to create and use a class diagram and have a solid footing for understanding other structure diagrams when you encounter them.

## The basics

As mentioned earlier, the purpose of the class diagram is to show the types and relationships in a system. In most UML models these types include:

- a class
- an interface
- a data type
- a component.

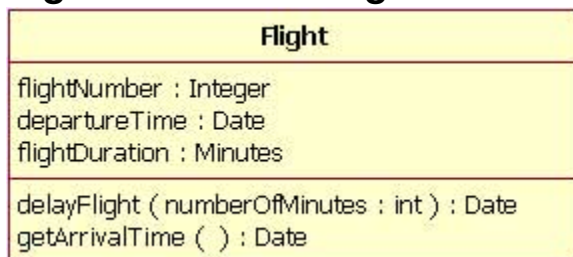
UML uses a special name for these types: "classifiers." Generally, you can use the term classifier, but technically a classifier is a more general term that refers to the other types of structure diagrams as well.

## Class name

The UML representation of a class is a rectangle containing three compartments. The top compartment contains the class name, the middle compartment contains the class attributes, and the bottom compartment contains the class operations.

shown in Figure 1. The top compartment shows the class's name. The middle compartment shows the class's attributes. The bottom compartment lists the class's operations. When creating a class diagram, you must use the top compartment, and the bottom two compartments are optional. (The bottom two would be unnecessary on a diagram depicting a higher-level relationship between classifiers.) Figure 1 shows the Flight class as a UML class. As we can see, the name is *Flight*, and in the middle compartment the class has three attributes: `flightNumber`, `departureTime`, and `flightDuration`. In the bottom compartment we see that the Flight class has two operations: `delayFlight` and `getArrivalTime`.

**Figure 1: Class diagram for the class Flight**



## Class attribute list

The attribute section of a class (the middle compartment) lists each of the attributes of the class. The attribute section is optional, but when used it contains each attribute in the following format. The line uses the following format:

name : attribute type
flightNumber : Integer

Continuing with our Flight class example, we can describe the class's attribute information, as shown in Table 1.

**Table 1: The Flight class's attribute names with their associated type**

Attribute Name	Attribute Type
flightNumber	Integer
departureTime	Date
flightDuration	Minutes

In business class diagrams, the attribute types usually correspond to units readers of the diagram (i.e., minutes, dollars, etc.). However, a class diagram generate code needs classes whose attribute types are limited to the type language, or types included in the model that will also be implemented in

Sometimes it is useful to show on a class diagram that a particular attribute example, in a banking account application a new bank account would start with a UML specification allows for the identification of default values in the attribute following notation:

```
name : attribute type = default value
```

For example:

```
balance : Dollars = 0
```

Showing a default value for attributes is optional; Figure 2 shows a Bank Account class called *balance*, which has a default value of 0.

**Figure 2: A Bank Account class diagram showing the balance attribute in dollars**



### Class operations list

The class's operations are documented in the third (lowest) compartment rectangle, which again is optional. Like the attributes, the operations of a class are documented in a specific format, with each operation on its own line. Operations are documented using the following notation:

```
name(parameter list) : type of value returned
```


The Flight class's operations are mapped in Table 2 below.

**Table 2: Flight class's operations mapped from Figure 3**

Operation Name	Parameters Return
delayFlight	<div> <div>Name</div> <div>numberOfMinutes</div> </div>
getArrivalTime	N/A

Figure 3 shows that the delayFlight operation has one input parameter — type Minutes. However, the delayFlight operation does not have a return value because I made a design decision not to have the delay operation should return the new arrival time, and if this were the case, it would appear as `delayFlight(numberOfMinutes : Minutes) : Date`. In UML, parameters, they are put inside the operation's parentheses; each parameter is in the form "name : parameter type".

**Figure 3: The Flight class operations parameters include the optional**



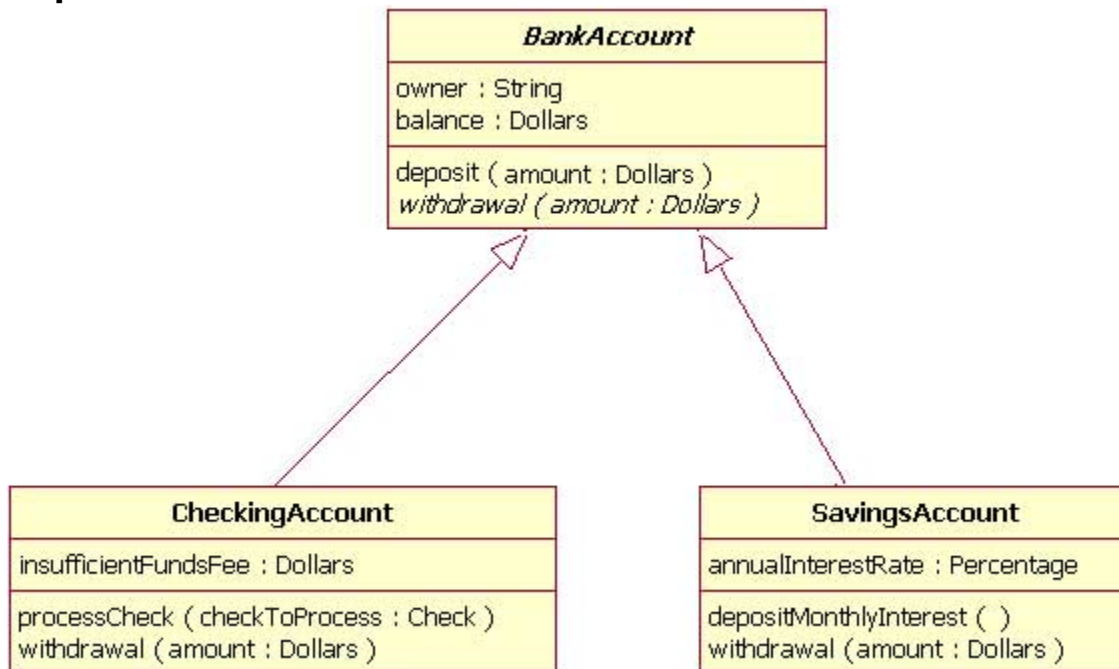
When documenting an operation's parameters, you may use an optional `in` or `out` keyword to indicate the parameter is input to, or output from, the operation. This optional indicator is shown in the operations compartment in Figure 3. Typically, these indicators are used in older programming languages such as Fortran, in which case they are helpful. However, in C++ and Java, all parameters are "in" parameters by default according to the UML specification, so most people will leave out the `in` keyword.

## Inheritance

A very important concept in object-oriented design, *inheritance*, refers to the relationship between a class and its superclass, where the subclass inherits the attributes and operations of the superclass.

class) to *inherit* the identical functionality of another class (super class), a its own. (In a very non-technical sense, imagine that I inherited my mothe in my family I'm the only one who plays electric guitar.) To model inheritan line is drawn from the child class (the class inheriting the behavior) with a triangle) pointing to the super class. Consider types of bank accounts: Fig CheckingAccount and SavingsAccount classes inherit from the BankAcco

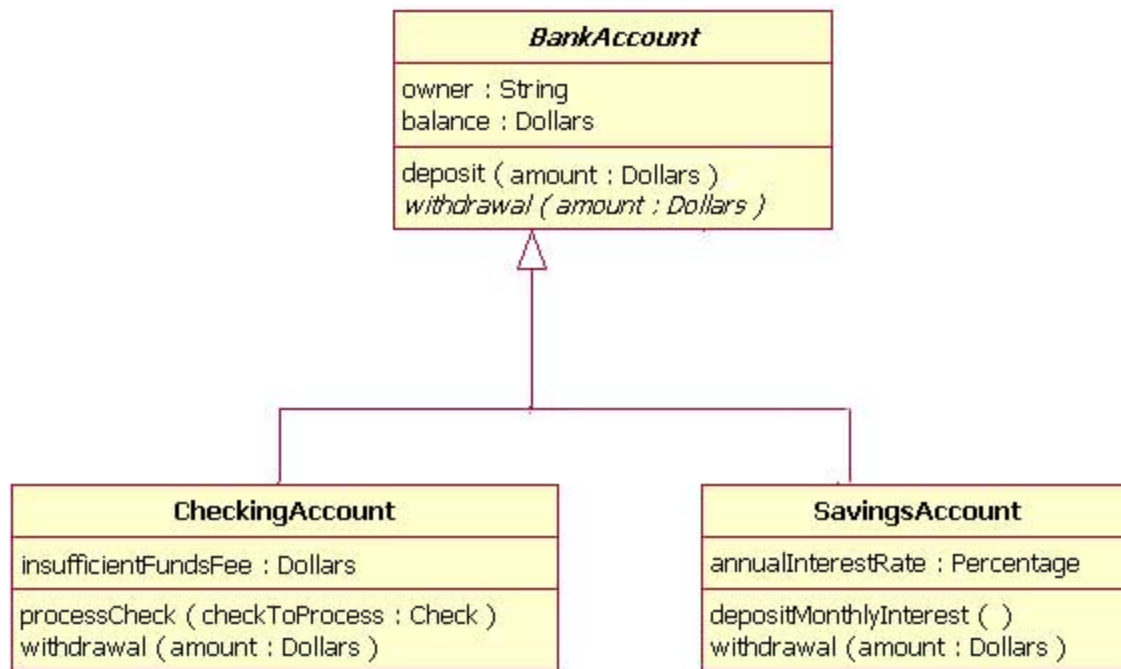
**Figure 4: Inheritance is indicated by a solid line with a closed, unfille super class**



In Figure 4, the inheritance relationship is drawn with separate lines for ea method used in IBM Rational Rose and IBM Rational XDE. However, ther inheritance called *tree notation*. You can use tree notation when there are Figure 4, except that the inheritance lines merge together like a tree bran the same inheritance shown in Figure 4, but this time using tree notation.

**Figure 5: An example of inheritance using tree notation**





## Abstract classes and operations

The observant reader will notice that the diagrams in Figures 4 and 5 use **BankAccount** class name and `withdrawal` operation. This indicates that **BankAccount** is an abstract class and the `withdrawal` method is an abstract operation. In other words, **BankAccount** provides the abstract operation signature of `withdrawal` and the two child classes, **CheckingAccount** and **SavingsAccount**, each implement their own version of that operation.

However, super classes (parent classes) do not have to be abstract class to be a super class.

## Associations

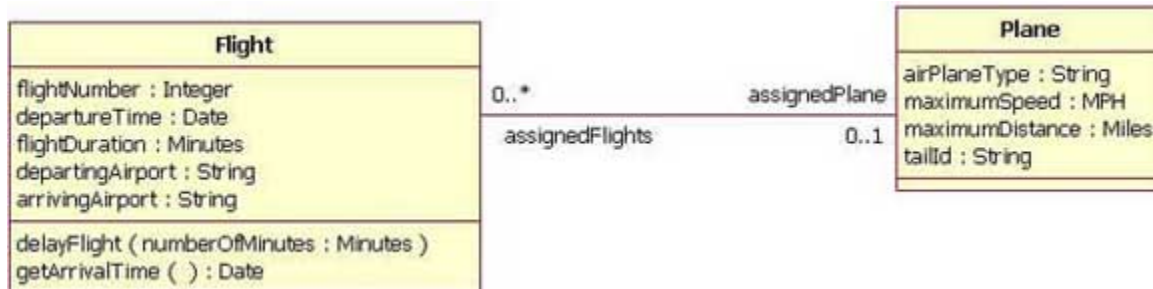
When you model a system, certain objects will be related to each other, and these relationships themselves need to be modeled for clarity. There are five types of associations — bi-directional and uni-directional associations — in this section, and I will discuss the other three association types in the *Beyond the basics* section. Please note that a detailed discussion of each type of association is beyond the scope of this article. Instead, I will focus on the first association type and show how the association is drawn on a class diagram.

### Bi-directional (standard) association



An association is a linkage between two classes. Associations are always this means that both classes are aware of each other and their relationship as some other type. Going back to our Flight example, Figure 6 shows an association between the Flight class and the Plane class.

**Figure 6: An example of a bi-directional association between a Flight**



A bi-directional association is indicated by a solid line between the two classes. You place a role name and a multiplicity value. Figure 6 shows that the **Flight** class knows about this association. The **Plane** class knows about this association because the role name next to the **Plane** class says so. The **Plane** class of `0..1` means that when an instance of a **Flight** exists, it can have one or no **Planes** associated with it (i.e., maybe a plane assigned). Figure 6 also shows that a **Plane** knows about its association with **Flight** because the **Flight** class takes on the role of "assignedFlights"; the diagram shows that a **Plane** instance can be associated either with no flights (e.g., it's a brand new plane) or with an infinite number of flights (e.g., the plane has been in commission for the last 10 years). For those wondering what the potential multiplicity values are for the ends of an association, Table 3 lists some example multiplicity values along with their meanings.

**Table 3: Multiplicity values and their indicators**

Potential Multiplicity Values	
Indicator	Meaning
0..1	Zero or one
1	One only

Indicator	Meaning
0..*	Zero or more
*	Zero or more
1..*	One or more
3	Three only
0..5	Zero to Five
5..15	Five to Fifteen

## Uni-directional association

In a uni-directional association, two classes are related, but only one class exists. Figure 7 shows an example of an overdrawn accounts report with :

**Figure 7: An example of a uni-directional association: The Overdraw knows about the BankAccount class, but the BankAccount class does not know about the OverdrawnAccountsReport class.**



A uni-directional association is drawn as a solid line with an open arrowhead (or triangle, used to indicate inheritance) pointing to the known class. Like a bi-directional association, the uni-directional association only contains the value for the known class. In our example in Figure 7, the **OverdrawnAccountsReport** class knows about the **BankAccount** class, and the **BankAccount** class plays the role of "overdrawnAccounts". In a standard association, the **BankAccount** class has no idea that it is associated with the **OverdrawnAccountsReport** class. [Note: It may seem strange that the **BankAccount** class does not know about the **OverdrawnAccountsReport** class. This modeling allows report classes to exist without knowing the business classes they report, but the business classes do not know they are being reported.

coupling of the objects and therefore makes the system more adaptive to

## Packages

Inevitably, if you are modeling a large system or a large area of a business classifiers in your model. Managing all the classes can be a daunting task organizing element called a *package*. Packages enable modelers to organize namespaces, which is sort of like folders in a filing system. Dividing a system makes the system easier to understand, especially if each package represents a system. [Note: Packages are great for organizing your model's classes, but that your class diagrams are supposed to easily communicate information modeled. In cases where your packages have lots of classes, it is better to use class diagrams instead of just producing one large class diagram.]

There are two ways of drawing packages on diagrams. There is no rule for use, except to use your personal judgement regarding which is easiest to use when drawing. Both ways begin with a large rectangle with a smaller rectangle in the top-left corner, as seen in Figure 8. But the modeler must decide how the package is drawn as follows:

If the modeler decides to show the package's members within the large rectangle, the members need to be placed within the rectangle. [Note: It's important to show those members," I mean only the classes that the current diagram is going to show. A package with contents does not need to show all its contents; it can show some elements according to some criterion, which is not necessarily all the package's contents. The package's name needs to be placed in the package's smaller rectangle (Figure 8).

If the modeler decides to show the package's members outside the large rectangle, the members that will be shown on the diagram need to be placed outside the rectangle. If a classifier belongs to the package, a line is drawn from each classifier to a circle that is attached to the package (Figure 9).

**Figure 8: An example package element that shows its members inside**

## boundaries

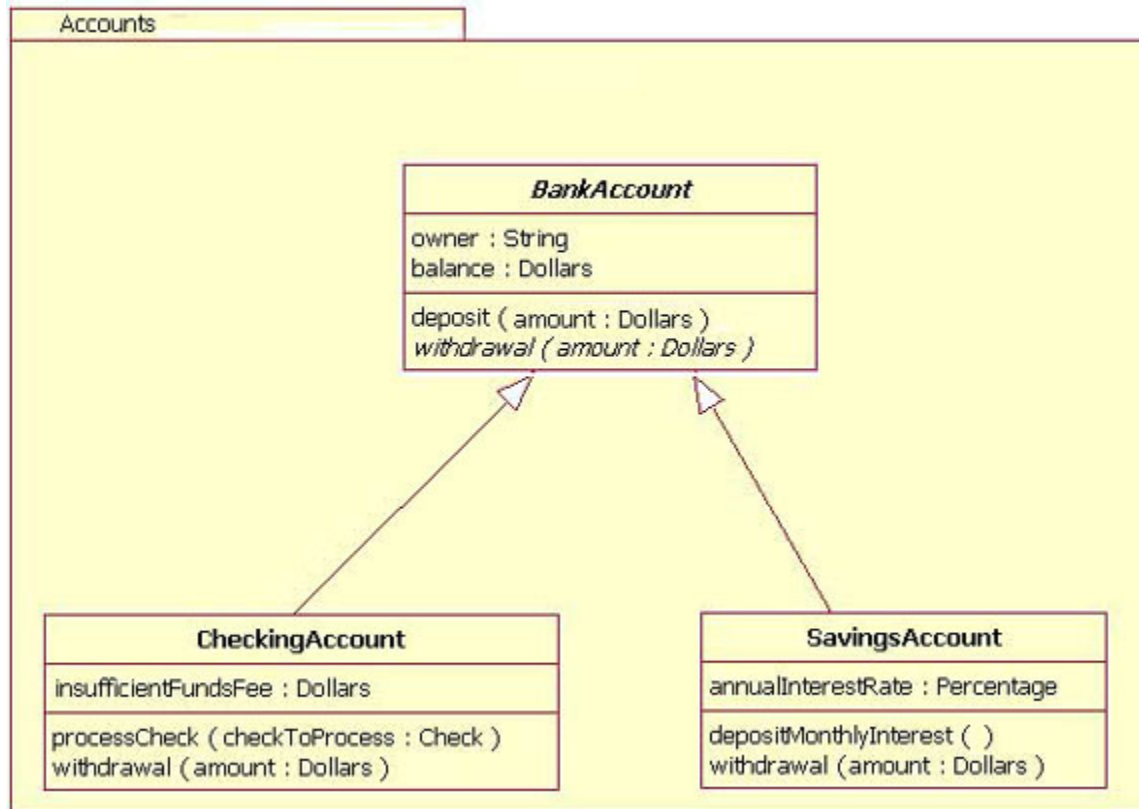
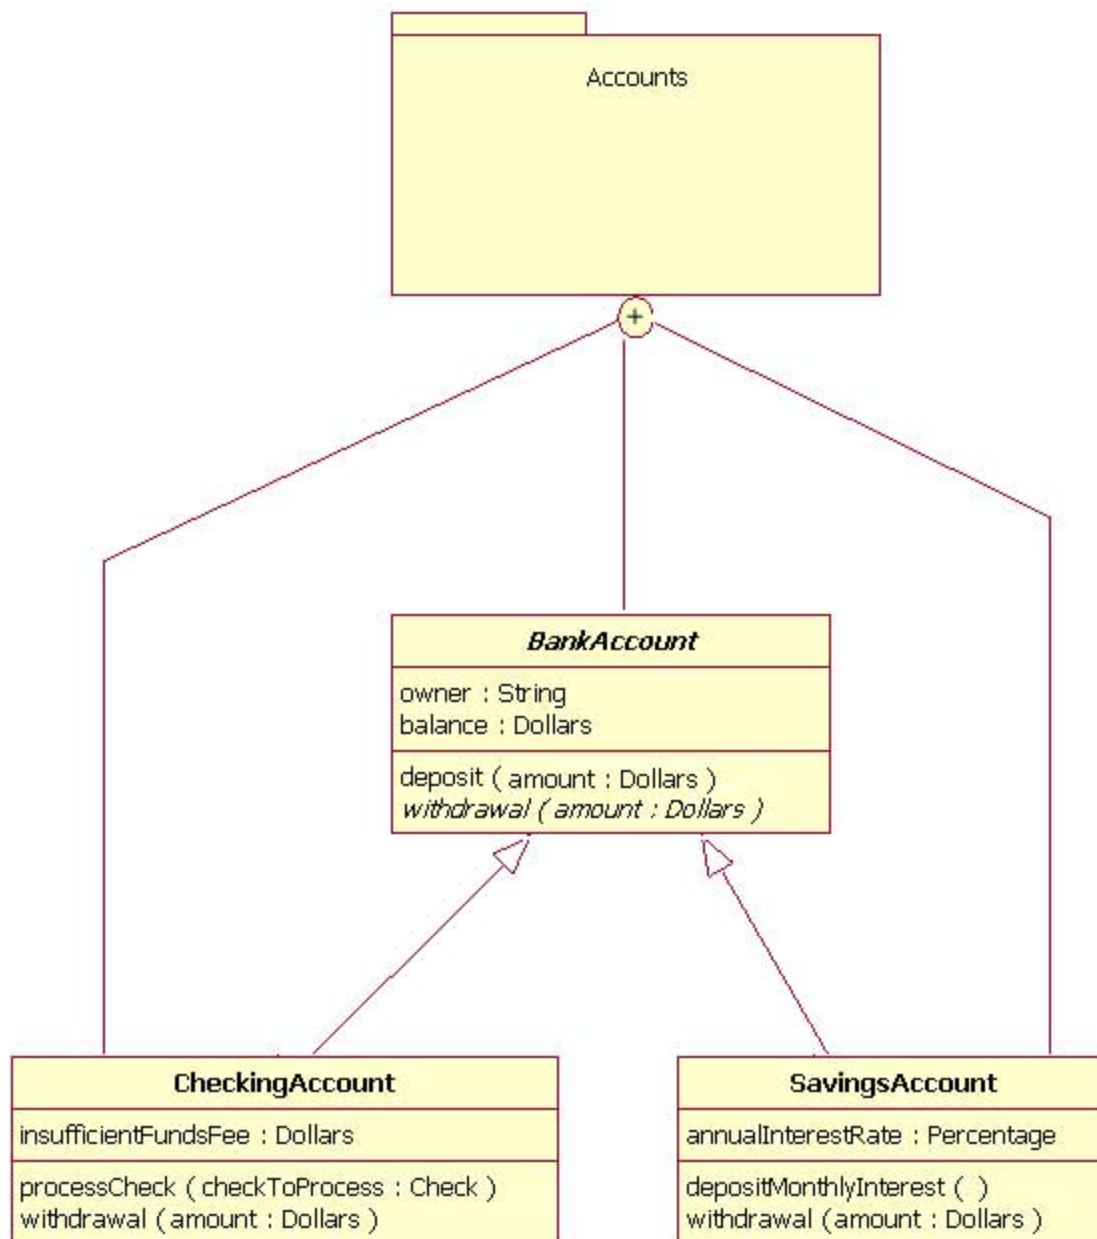


Figure 9: An example package element showing its membership via



## Importance of understanding the basics

It is more important than ever in UML 2 to understand the basics of the class diagram. The class diagram provides the basic building blocks for all other structure component or object diagrams (just to name a few).

## Beyond the basics

At this point, I have covered the basics of the class diagram, but do not st

sections, I will address more important aspects of the class diagram that y include interfaces, the three remaining types of associations, visibility, and specification.

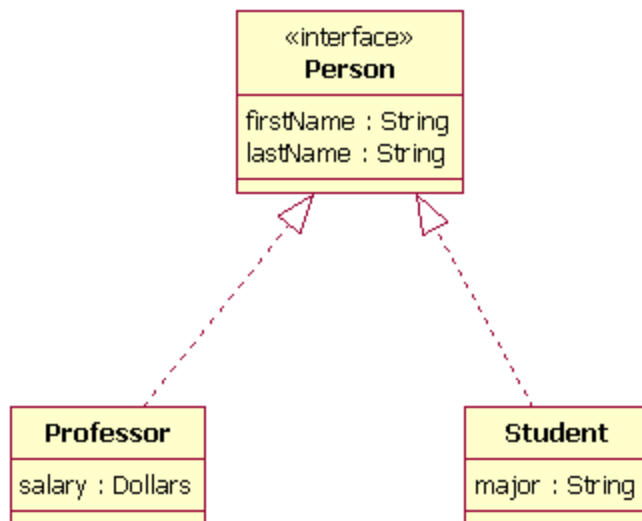
## Interfaces

Earlier in this article, I suggested that you think of *classifiers* simply as cla more general concept, which includes data types and interfaces.

A complete discussion of when and how to use data types and interfaces structure diagrams is beyond the scope of this article. So why do I mentio here? There are times when you might want to model these classifier type is important to use the proper notation in doing so, or at least be aware of these classifiers incorrectly will likely confuse readers of your structure dia will probably not meet requirements.

A class and an interface differ: A class can have an actual instance of its t have at least one class to implement it. In UML 2, an interface is consider class modeling element. Therefore, an interface is drawn just like a class, rectangle also has the text "«interface»", as shown in Figure 10. [Note: W is completely within UML specification to put «class» in the top compartm would with «interface»; however, the UML specification says that placing i compartment is optional, and it should be assumed if «class» is not displa

**Figure 10: Example of a class diagram in which the Professor and St Person interface**



In the diagram shown in Figure 10, both the Professor and Student classes implement the `Teachable` interface and do not inherit from it. We know this for two reasons: 1) The `Teachable` object is an interface — it has the "«interface»" text in the object's name area, and we know that Professor and Student objects are *class* objects because they are labeled according to that classification (there is no additional classification text in their name area). 2) We know that Professor and Student classes implement the `Teachable` interface, as shown here, because the line with the arrow is dotted and not solid. As shown in Figure 9, a line with a closed, unfilled arrow means realization (or implementation); as we saw in Figure 8, a line with a closed, unfilled arrow means inheritance.

## More associations

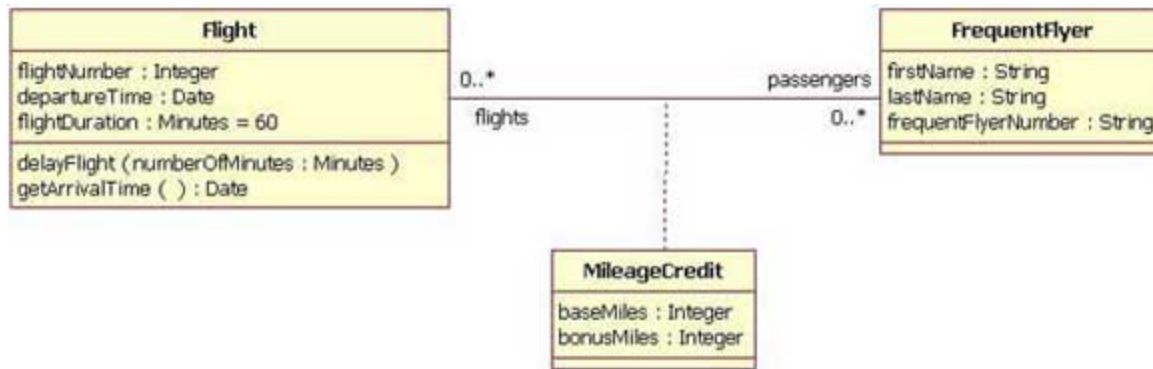
Above, I discussed bi-directional and uni-directional associations. Now I v types of associations.

## Association class

In modeling an association, there are times when you need to include an association class. An association class is a class that contains valuable information about the relationship. For this you would use an association class. An association class is represented like a normal class. An association line between the primary classes intersects a dotted line connecting the association class. Figure 11 shows an association class for our airline industry example.

### Figure 11: Adding the association class MileageCredit





In the class diagram shown in Figure 11, the association between the **Flight** class results in an association class called **MileageCredit**. This means that if a class is associated with an instance of a **FrequentFlyer** class, there will always be a **MileageCredit** class.

## Aggregation

Aggregation is a special type of association used to model a "whole to its parts" relationship. In aggregation relationships, the lifecycle of a *part* class is independent from the lifecycle of the *whole* class.

For example, we can think of *Car* as a whole entity and *Car Wheel* as a part. A wheel can be created weeks ahead of time, and it can sit in a warehouse before assembly. In this example, the *Wheel* class's instance clearly lives independently of the *Car* class's instance. However, there are times when the *part* class's lifecycle is *not* independent of the *whole* class — this is called composition aggregation. Consider, for example, a company and its departments. Both *Company* and *Departments* are modeled as classes. A department cannot exist before a company exists. Here the *Department* class's instance is dependent on the existence of the *Company* class's instance.

Let's explore basic aggregation and composition aggregation further.

### Basic aggregation

An association with an aggregation relationship indicates that one class is the whole and the other is a part. In an aggregation relationship, the child class instance can outlive its parent class instance. To represent an aggregation relationship, you draw a solid line from the parent class to the part class, with an open circle at the part class end.

shape on the parent class's association end. Figure 12 shows an example between a Car and a Wheel.

**Figure 12: Example of an aggregation association**



## Composition aggregation

The composition aggregation relationship is just another form of the aggregation relationship. The child class's instance lifecycle is dependent on the parent class's instance lifecycle. Figure 13 shows a composition relationship between a Company class and a Department class. A composition relationship is drawn like the aggregation relationship, but the diamond is filled.

**Figure 13: Example of a composition relationship**




In the relationship modeled in Figure 13, a Company class instance will always have a Department class instance. Because the relationship is a composition relationship, when the Company instance is removed/destroyed, the Department instance is automatically removed/destroyed. Another important feature of composition aggregation is that the part class's instance lifecycle is dependent on the parent class's instance lifecycle (e.g. the Company class in our example).

## Reflexive associations

We have now discussed all the association types. As you may have noticed, a class can have a relationship between two different classes. However, a class can also have a relationship with itself using a reflexive association. This may not make sense at first, but remember that classes are abstractions. Figure 14 shows how an Employee class could be related to itself using a reflexive association. For example, an Employee class could have a manager/manages role. When a class is associated to itself, this does not mean that the class is related to itself, but that an instance of the class is related to another instance of the class.

**Figure 14: Example of a reflexive association relationship**



The relationship drawn in Figure 14 means that an instance of Employee Employee instance. However, because the relationship role of "manages" Employee might not have any other Employees to manage.

## Visibility

In object-oriented design, there is a notation of visibility for attributes and types of visibility: public, protected, private, and package.

The UML specification does not require attributes and operations visibility diagram, but it does require that it be defined for each attribute or operation in a class diagram, you place the visibility mark in front of the attribute's or operation's name. The UML specification specifies four visibility types, an actual programming language may add a visibility type that the language does not support the UML-defined visibilities. Table 4 displays the different marks for the four visibility types.

**Table 4: Marks for UML-supported visibility types**

Mark	Visibility type
+	Public
#	Protected
-	Private
~	Package

Now, let's look at a class that shows the visibility types indicated for its attributes and operations. In Figure 15, all the attributes and operations are public, with the exception of the updateBalance operation. The updateBalance operation is protected.

**Figure 15: A BankAccount class that shows the visibility of its attributes and operations**



## UML 2 additions

Now that we have covered the basics and the advanced topics, we will co added to the class diagram from UML 1.x.

### Instances

When modeling a system's structure it is sometimes useful to show exam model this, UML 2 provides the *instance specification* element, which sho example (or real) instances in the system.

The notation of an instance is the same as a class, but instead of the top class's name, the name is an underlined concatenation of:


Instance Name : Class Name
----------------------------

For example:

Donald : Person
-----------------

Because the purpose of showing instances is to show interesting or relev necessary to include in your model the entire instance's attributes and op appropriate to show only the attributes and their values that are interesting

**Figure 16: An example instance of a Plane class (only the interesting**



However, merely showing some instances without their relationship is not allows for the modeling of the relationships/associations at the instance le

drawing associations are the same as for normal class relationships, although there is an additional requirement when modeling the associations. The association must match the class diagram's relationships and therefore the association must be drawn in the same way as the class diagram. An example of this is shown in Figure 17. In this example, the association is between instances of the class diagram found in Figure 6.

**Figure 17: An example of Figure 6 using instances instead of classes**

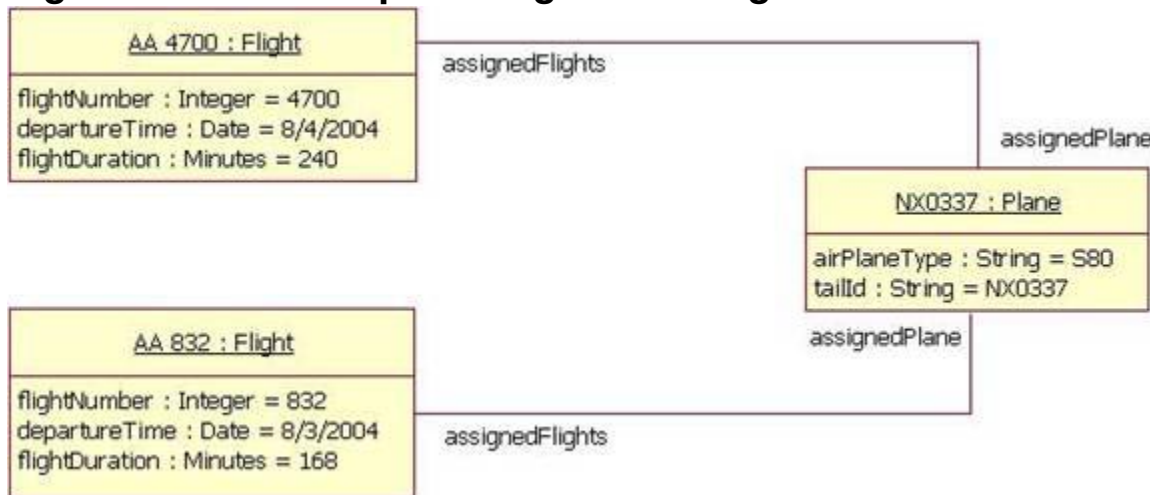


Figure 17 has two instances of the Flight class because the class diagram between the Plane class and the Flight class is *zero-to-many*. Therefore, Flight instances that the NX0337 Plane instance is related to.

## Roles

Modeling the instances of classes is sometimes more detailed than one might simply want to model a class's relationship at a more generic level. In such cases, the *role* notation is used. The role notation is very similar to the instances notation. To use it, you place the class's role name and class name inside a box as with the instances notation. In this case, you do not underline the words. Figure 18 shows an example of the class described by the diagram at Figure 14. In Figure 18, we can tell, even if the class is related to itself, that the relationship is really between an Employee playing the role of team member.

**Figure 18: A class diagram showing the class in Figure 14 in its different roles**



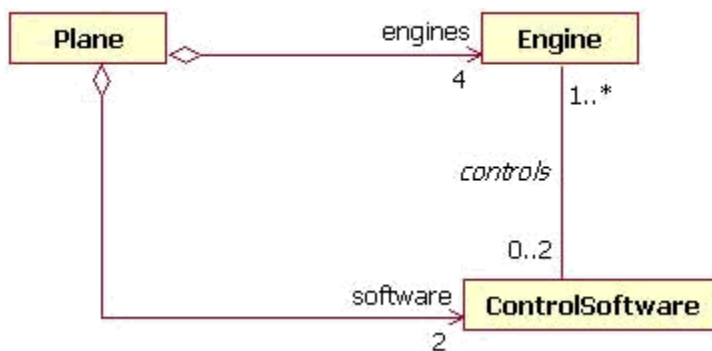
Note that you cannot model a class's role on a plain class diagram, even if it appears that you can. In order to use the role notation you will need to use the notation discussed next.

## Internal Structures

One of the more useful features of UML 2 structure diagrams is the new internal structure notation, which allows you to show how a class or another classifier is internally composed. In UML 1.x, because the notation set limited you to showing only the aggregation relationship, Now, in UML 2, the internal structure notation lets you more clearly show how parts are related to each other.

Let's look at an example. In Figure 18 we have a class diagram showing the relationships between four engines and two control software objects. What is missing from this diagram is information about how airplane parts are assembled. From the diagram in Figure 18, it is not clear if two control software objects control two engines each, or if one control software object controls two engines and the other controls one engine.

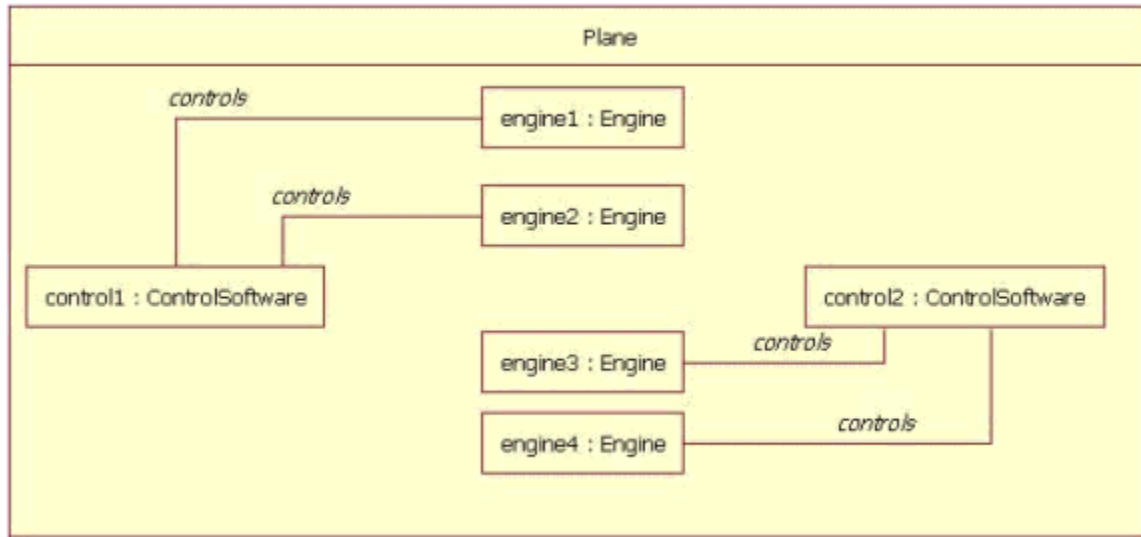
**Figure 19: A class diagram that only shows relationships between the**



Drawing a class's internal structure will improve this situation. You start by drawing the class's internal structure. The top compartment contains the class name, and the bottom compartment contains the class's internal structure, showing the parent class's part classes in their relationships.

each particular class relates to others in that role. Figure 19 shows the int notice how the internal structure clears up the confusion.

**Figure 20: An example internal structure of a Plane class**



In Figure 20 the Plane has two ControlSoftware objects and each one controls two Engine objects. ControlSoftware on the left side of the diagram (control1) controls engines 1 and 2 on the right side of the diagram (control2) controls engines 3 and 4.

## Conclusion

There are at least two important reasons for understanding the class diagram notation for other structure diagrams prescribed by UML. Developers will create them specially for them; but other team members will find them useful, class diagrams to model systems from the business perspective. As we continue this series on UML basics, other diagrams — including the activity, sequence, and state machine diagrams — refer to the classes modeled and documented on the class diagram.

Next in this series on UML basics: [The component diagram](#).

---

## Resources

D  
O



## Learn

Understand more about how the new functionality of the UML Modeler component common to both IBM Rational Software Architect Standard Edition Version 7.5 and IBM Rational Software Architect for WebSphere Software Version 7.5 by reading this article [Using the new features of UML Modeler in IBM Rational Software Architect Version 7.5](#).

Understand more about the new functionality of Rational Software Architect for WebSphere Software Version 7.5 by reading this article [Overview of Rational Software Architect for WebSphere Software Version 7.5](#).

Learn about [Rational Tau](#).

Learn about other applications in the [IBM Rational Software Delivery Platform](#), including collaboration tools for parallel development and geographically dispersed teams, plus specialized software for architecture management, asset management, change and release management, integrated requirements management, process and portfolio management, and quality management. You can find product manuals, installation guides, and other documentation in the [IBM Rational Online Documentation Center](#).

Visit the [Rational software area on developerWorks](#) for technical resources and best practices for Rational Software Delivery Platform products.

Explore [Rational computer-based, Web-based, and instructor-led online courses](#). Hone your skills and learn more about Rational tools with these courses, which range from introductory to advanced. The courses on this catalog are available for purchase through computer-based training or Web-based training. Some of the "Getting Started" courses are available free of charge.

Subscribe to the [IBM developerWorks newsletter](#), a weekly update on the best of developerWorks tutorials, articles, downloads, community activities, webcasts and events.

Browse the [technology bookstore](#) for books on these and other technical

topics.

## Get products and technologies

Download a [trial version of Rational Software Modeler](#).

Download a [trial version of Rational Software Architect](#) standard edition.

Download [Rational Software Architect for Websphere Software](#).

Download [Other IBM product evaluation versions](#) and get your hands on application development tools and middleware products from DB2®, Lotus®, Rational®, Tivoli®, and WebSphere®.

## Discuss

Check out [developerWorks blogs](#) and get involved in the [developerWorks community](#).

Join the [Development Tools forum](#) on developerWorks to discuss Rational Application Developer, Rational Software Architect, and Rational Software Modeler.

Join the [Rational Tau forum](#) on developerWorks.